

Assignment 2 _Group 6

Baochen Hu & Ning Zhang

Kafka word count:

Python:

1. Start Zookeeper

bin/zookeeper-server-start.sh config/zookeeper.properties

2. configuration

/config/server.properties

host.name = localhost

advertised.host.name = localhost

3. Start Kafka Broker

bin/kafka-server-start.sh config/server.properties

4. create topic

bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor

1 -- partitions 1 --topic test

bin/kafka-topics.sh --list --zookeeper localhost:2181

5. Start a Producer

bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test

6. Start a consumer

bin/kafka-console-consumer.sh --zookeeper localhost:2181 --topic test --from-beginning

7. download a reliable jar, and put it in /external/kafka-assembly

wget http://search.maven.org/remotecontent?filepath=org/apache/spark/spark-streaming-kafka-assembly_2.10/1.4.0/spark-streaming-kafka-assembly_2.10-1.4.0.jar

6.

bin/spark-submit --jars

external/kafka-assembly/spark-streaming-kafka-assembly_2.10-1.4.0.jar

examples/src/main/python/streaming/kafka_wordcount.py localhost:2181 test

Scala:

```
bin/run-example org.apache.spark.examples.streaming.KafkaWordCount localhost:2181
test-consumer-group test 1
```

```
-----
Time: 2015-08-08 04:36:14
-----
```

```
(u'hello', 3)
()
```

Flume word count

1. apt-get <http://apache.arvix.com/flume/1.6.0/apache-flume-1.6.0-bin.tar.gz>

2. tar -xvz apache-flume-1.6.0-bin.tar.gz

3. add config.properties file.

```
a1.sources = r1
```

```
a1.sinks = k1
```

```
a1.channels = c1
```

```
# Describe/configure the source
```

```
a1.sources.r1.type = netcat
```

```
a1.sources.r1.bind = 127.0.0.1
```

```
a1.sources.r1.port = 44444
```

```
# Describe the sink
```

```
#a1.sinks = avroSink
```

```
a1.sinks.k1.type = avro
```

```
a1.sinks.k1.channel = memoryChannel
```

```
a1.sinks.k1.hostname = 127.0.0.1
```

```
a1.sinks.k1.port = 8989
```

```
#a1.sinks.k1.type = logger
```

Use a channel which buffers events in memory

a1.channels.c1.type = memory

a1.channels.c1.capacity = 5

a1.channels.c1.transactionCapacity = 5

Bind the source and sink to the channel

a1.sources.r1.channels = c1

a1.sinks.k1.channel = c1

4.bin/flume-ng agent --conf conf --conf-file conf/flume-spark-integration.properties
--name a1 -Dflume.root.logger=INFO,console

5. open another terminal: telnet localhost 44444

6.open another terminal: bin/run-example
org.apache.spark.examples.streaming.FlumeEventCount 127.0.0.1 8989

NOTE: This should be ran first.

```
-----  
Tip root@li796-54: ~/apache-flume-1.6.0-bin — ssh — 159x20  
-----
```

Received 1 flume events.

```
-----  
Time: 1438999076000 ms  
-----
```

Received 0 flume events.

```
-----  
Time: 1438999078000 ms  
-----
```

Received 0 flume events.

```
-----  
Time: 1438999080000 ms  
-----
```

Received 0 flume events.

```

Trying ::1...
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
hello
OK
ni
OK
Connection closed by foreign host.
root@li796-54:~# telnet localhost 44444
Trying ::1...
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
hello
OK
hhllo
OK
^CConnection closed by foreign host.

```

HDFS word count:

hadoop fs -ls /user

hadoop fs -mkdir /user/ningzhang/sparkStreaming

scala:

Counts words in new text files created in the given directory

* Usage: HdfsWordCount <directory>

* <directory> is the directory that Spark Streaming will use to find and read new text files.

*

* To run this on your local machine on directory `localdir`, run this example

* \$ bin/run-example \

* org.apache.spark.examples.streaming.HdfsWordCount localdir

*

* Then create a text file in `localdir` and the words in the file will get counted.

*/

bin/run-example org.apache.spark.examples.streaming.HdfsWordCount

hdfs://ec2-52-2-200-145.compute-1.amazonaws.com/data/

hdfs dfs -copyFromLocal /home/hadoop/txtFile/hello.txt

hdfs://ec2-52-2-200-145.compute-1.amazonaws.com/data/hello1

Time: 1438993764000 ms

(hi,3)
(hello,2)

MQTT word count:

Install mosquitto:

1. <http://mosquitto.org/download/>

For Centos

Add the CentOS mosquitto repository to YUM's list of repositories

```
$ cd /etc/yum/yum.repos.d
```

```
$ sudo
```

```
wget http://download.opensuse.org/repositories/home:/oojah:/mqtt/CentOS_CentO...
```

```
$ sudo yum update
```

```
$ sudo yum install mosquitto
```

As of writing, no init.d script exists for the CentOS distribution of mosquitto. However, it is a simple enough matter to set it running as a daemon, you'll just need to restart it yourself whenever your machine gets restarted

```
1. $ sudo su
```

```
2. $ /usr/sbin/mosquitto -d -c /etc/mosquitto/mosquitto.conf > /var/log/mosquitto.log  
2>&1
```

```
$mosquitto
```

To run this example locally, you may run publisher as

```
bin/run-example
```

```
org.apache.spark.examples.streaming.MQTTPublisher tcp://localhost:1883 foo
```

```
bin/run-example
```

```
org.apache.spark.examples.streaming.MQTTWordCount tcp://localhost:1883 foo
```

TwitterPopularTags:

scala:

```
1.
```

```
./bin/run-example org.apache.spark.examples.streaming.TwitterPopularTags
```

```
fLFWZflOEZbUhPjhtfYvtmycy
```

```
iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde
```

```
2612760793-cSRGyHgdIByaMeQZojisB6HYCJaPw9rgYsokbqjL
```

```
CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil
```

<http://spark.apache.org/docs/latest/quick-start.html>

```
1. install sbt
```

```
sudo yum update
```

```
curl https://bintray.com/sbt/rpm/rpm | sudo tee
/etc/yum.repos.d/bintray-sbt-rpm.repo
sudo yum install sbt
```

2. dependency: simple.sbt

```
name := "Simple Project"
```

```
version := "1.0"
```

```
scalaVersion := "2.10.4"
```

```
libraryDependencies ++= Seq(
```

```
  "org.apache.spark" %% "spark-core" % "1.4.1",
```

```
  "org.apache.spark" %% "spark-streaming" % "1.4.1",
```

```
  "org.apache.spark" %% "spark-streaming-twitter" % "1.4.1"
```

```
)
```

```
vi TwitterPopularTags.scala
```

```
/* StreamingExamples.setStreamingLogLevels() */
```

3. sbt package

4. go to twitter developer to get token

<https://apps.twitter.com/app/new>

Consumer Key (API Key) fLFWZfIOEZbUhPjhtfYvtmycy

Consumer Secret (API

Secret) iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde

Access Token 2612760793-cSRGyHgdIByaMeQZojSB6HYCJaPw9rgYsokbqjL

Access Token Secret CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil

5.

```
./../../usr/lib/spark/bin/spark-submit
```

```
twitterPopularTags/target/scala-2.10/simple-project_2.10-1.0.jar
```

```
fLFWZfIOEZbUhPjhtfYvtmycy
```

```
iPGu1R7NGTt8SQkRpmFC1OfSs3VOC6vJwhzBP7yhABYTnzNvde
```

```
2612760793-cSRGyHgdIByaMeQZojSB6HYCJaPw9rgYsokbqjL
```

```
CACXdrOf7XDnWOUrc94fdawucDs56Ylz7GZRtRpFRDnil
```

```

Popular topics in last 60 seconds (50 total):
#AlDubWeBelongTogether (2 tweets)
#flycosta (1 tweets)
#mizzoudg (1 tweets)
#騒音トラブル
http://t.co/WnXFwef8GC (1 tweets)
#sistas (1 tweets)
#online... (1 tweets)
#Internet (1 tweets)
#里親募集 (1 tweets)
#Follow (1 tweets)
#football (1 tweets)
15/08/08 05:01:24 WARN BlockManager: Block input-0-1439010084200 replicated to only 0 peer(s) instead of 1 peers

```

```

Popular topics in last 10 seconds (50 total):
#AlDubWeBelongTogether (2 tweets)
#flycosta (1 tweets)
#mizzoudg (1 tweets)
#騒音トラブル
http://t.co/WnXFwef8GC (1 tweets)
#sistas (1 tweets)
#online... (1 tweets)
#Internet (1 tweets)
#里親募集 (1 tweets)
#Follow (1 tweets)
#football (1 tweets)

```

ZeroMQ word count:

Install ZeroMQ:

<https://www.digitalocean.com/community/tutorials/how-to-install-zeromq-from-source-on-a-centos-6-x64-vps>

wget <http://download.zeromq.org/zeromq-2.1.10.tar.gz>

2.0.10

```

bin/run-example org.apache.spark.examples.streaming.SimpleZeroMQPublisher tcp://
127.0.1.1:1234 foo.bar
bin/run-example org.apache.spark.examples.streaming.ZeroMQWordCount tcp://127.
0.1.1:1234 foo

```

```

15/08/07 23:55:58 WARN BlockManager: Block input-0-1438991757105 replicated to only 0 peer(s) instead of 1 peers
15/08/07 23:55:59 WARN BlockManager: Block input-0-1438991757106 replicated to only 0 peer(s) instead of 1 peers
-----
Time: 1438991760000 ms
-----
(count,2)
(words,2)
(may,2)

```